

Persona-Aware Recommendation for Nigerian E-Commerce: An Agentic Pipeline with an 8B Open-Weight Fine-Tune as Re-Ranker

Ashinze Emmanuel
International University of Applied
Sciences, Germany

Franca Uvere
University of Lagos
Nigeria

Esther Oyeneke
MIVA Open University
Nigeria

Abstract

We address the Recommendation track of the DSN X BCT LLM Agent Challenge such that given a Nigerian-context persona and a candidate set of products, rank those products in order of likely user preference, handling the three sub-scenarios called out in the brief; **cold-start** (no purchase history), **cross-domain** (recommending across product categories the user has never bought in), and **multi-turn** (conversational refinement of recommendations). We present a six-stage agentic pipeline built on (1) a structured cognitive persona representation shared with our Task-A user-modeling work, (2) asymmetric dense retrieval over 6 657 real Jumia products using llama-text-embed-v2 hosted on Pinecone, (3) a *Cohere rerank-v3.5 cross-encoder pre-rerank* that narrows 30 candidates to 15 by persona-flavored query before any LLM is called, (4) a configurable persona-aware LLM re-ranker (any of 11 backbones), and (5) Maximal Marginal Relevance for category diversity. The agent surfaces its own scenario detection (cold_start, cross_domain, multi_turn flags) and a narrative reasoning trace explaining each ranking decision.

The headline empirical result is that our 8 B parameter open-weight fine-tune, **NaijaReviewer-8B** (Llama 3.1 8B QLoRA-trained on Nigerian product reviews, originally for Task A), is also the *strongest re-ranker* we tested for Task B in *both* pipeline configurations, with and without the Stage 2.5 Cohere cross-encoder pre-rerank enabled. On a 17–18 persona-scenario held-out split, Cohere-OFF: **NDCG@10 0.588** (vs Claude 0.433, GPT-OSS-120B 0.441, Llama-3.3-70B 0.425, Qwen-2.5-72B 0.404); Cohere-ON: **NDCG@10 0.572** (vs Llama-3.3-70B 0.477, Qwen-2.5-72B 0.461, Claude 0.430, GPT-OSS-120B 0.366). NaijaReviewer-8B wins on NDCG@10 across both configurations at 9–15× fewer parameters than the next-best baseline, and at zero per-call API cost. Cohere narrowing lifts NaijaReviewer-8B’s HR@1 by +19% (0.444 → 0.529), it gets the top-n pick right more often when the input pool is tightened. We audit the result via five controlled live rerank calls and a fallback-event count over the full eval: zero fallback events fired, confirming the win is the model’s actual re-ranking logic and not a pre-rank backoff.

1 Introduction

A Lagos-based Bolt driver shopping on Jumia is not the same recommendation target as a Kano teacher, even if both have purchased a 4-star phone case. Their aspect priorities differ (durability vs aesthetics), their budgets differ, the brands they trust differ, and the products that count as “culturally appropriate” for their household differ. Frontier recommendation systems trained on Amazon / Western e-commerce data do not see these distinctions; they recommend

by global review-text similarity, which collapses Nigerian shoppers into a generic “mobile-phone-buyer” archetype.

This paper addresses **Task B** of the DSN X BCT LLM Agent Challenge, given a persona, return a ranked list of products from a candidate set, handling cold-start, cross-domain, and multi-turn sub-scenarios. A companion paper addresses Task A (User Modeling / review generation) using the same persona substrate.

Contributions.

- A *six-stage agentic recommendation pipeline* with explicit handlers for the three Task B sub-scenarios (cold-start popularity branching, cross-domain de-duplicated retrieval + MMR diversity bias, multi-turn constraint extraction). The agent surfaces its own scenario detection flags and a narrative reasoning_trace explaining each ranking decision, the “agentic workflow that reasons before recommending” language the brief calls for.
- A *persona-aware retrieval architecture* built on llama-text-embed-v2 (1024-dim, Llama-3.2-1B encoder) hosted on Pinecone with the same asymmetric passage / query encoding architecture used by industrial career / skills-extraction systems, indexing 6 657 real Jumia products.
- A *Cohere rerank-v3.5 cross-encoder pre-rerank* stage that narrows the post-retrieval candidate pool 30 → 15 by persona-flavored query in ~200–600 ms before the (slow, expensive) LLM rerank stage. It auto-skips gracefully when the pre-rerank is unavailable or the pool is already small enough.
- An *11-backbone re-ranker comparison* on a 17–18 persona-scenario held-out split. The 8 B *NaijaReviewer-8B* open-weight fine-tune is the top re-ranker on every metric, beating Claude Sonnet 4, GPT-OSS-120B, Llama-3.3-70B, and Qwen-2.5-72B at 9–15× fewer parameters and zero per-call API cost.
- A *robust JSON-following pipeline*: a custom fence-stripping, balanced-brace, and truncated-JSON-repair parser that lets smaller and less-aligned models participate in re-ranking without fallback-to-pre-rank events. Verified across 18 scenarios: zero fallbacks fired.
- An *open data release*: 24 hand-curated Nigerian personas spanning 6 geopolitical zones × 4 age bands × 4 register tiers × 3 religions, plus a filtered 6 657-product real-Jumia catalogue. All released CC-BY-4.0.
- A *reproducible open-source system*: FastAPI service (POST /recommend, POST /chat), React/Vite frontend, eval harness, everything one command from a fresh clone.

2 Related Work

Our pipeline sits at the intersection of four research threads: dense retrieval, cross-encoder and LLM re-ranking, persona-aware recommendation, and low-resource African NLP. Table 1 positions the present work against the closest prior systems.

Dense retrieval and cross-encoder re-ranking. Dense passage retrieval with a dual-encoder architecture [4, 9] is now the standard first stage in modern retrieval pipelines. A slower cross-encoder, exemplified by monoBERT-style re-ranking [8] and the Cohere rerank-v3.5 model we use, typically refines the top candidates. Our pipeline keeps both, then adds an LLM re-rank as a third tier, mirroring the cost-and-quality trade-off of modern enterprise search. We use llama-text-embed-v2 with separate passage and query encoders, an asymmetric scheme also used by industrial skills-extraction systems [5].

LLMs as re-rankers. Recent work demonstrates that frontier LLMs can match or exceed classical cross-encoders for relevance re-ranking [11, 14], and that listwise prompting helps calibrate scores across candidates. Our addition is to condition the LLM re-rank on a structured cognitive persona (below) rather than only a query string, and to compare eleven backbones on the same candidate set, isolating the re-ranker as the variable.

Persona-aware and conversational recommendation. AgentCF++ [15] and RecAgent [13] represent users by concatenated rating histories and natural-language profiles. Conversational recommendation [12] extends this to multi-turn refinement of preferences. We extend both lines by replacing free-form text with a structured behavioural representation that is shared verbatim with the companion Task A paper, so the same persona JSON drives both review generation and re-ranking.

Cold-start. Cold-start recommendation has been studied since the early 2000s [10]. Our pipeline uses an explicit popularity / aspect-match / similarity branch when a persona has little or no history, surfaced through the agent’s reasoning trace.

Adapting open-weight LLMs. NaijaReviewer-8B is a QLoRA fine-tune of Llama 3.1 8B Instruct [7] using the LoRA / QLoRA recipe of Hu et al. [3] and Dettmers et al. [2]. Retrieval-augmented generation [6] is the architectural template our re-ranker follows: a learned retriever feeds a conditional language model.

African-language NLP and Nigerian context. A release of the largest publicly-available Nigerian product review corpus was found on GitHub and Huggingface; we filter it for the 6 657-product Jumia catalogue we index. The Masakhane community [1] has produced the canonical African-language pretraining checkpoints. We use these indirectly through the fine-tuned re-ranker, which is trained on ~20k Nigerian product reviews.

3 Method

3.1 Cognitive Persona Representation

A persona in our system is a 6-field JSON record: hedonic_utilitarian $\in [0, 1]$, communal_individual $\in [0, 1]$, intensity_calibration (a mapping from intensifier words to predicted star ratings, e.g. {“amazing”:4.8, “okay”:3.0}), aspect_priority (a probability vector over

{quality, value, delivery, packaging, seller, durability, design}), register_tier $\in \{\text{Nigerian Pidgin, code-mixed, Nigerian English, standard English}\}$, and review_anchors (a list of past reviews with their star rating, used for cold-start backoff). The same representation is used by the Task-A review-generation agent (companion paper); the persona is the unit of transfer between the two tasks. For Task B specifically, the agent uses aspect_priority as the primary persona-product fit signal and hedonic_utilitarian / communal_individual as tie-breakers in the re-rank stage.

3.2 Recommendation Pipeline

The agent runs six stages per request, each visible in the returned reasoning trace:

(0) Intent analysis. If a conversation_history is provided (multi-turn scenarios per the brief), a regex-grounded constraint extractor pulls budget, recipient, and category constraints. A flag set (cold_start, cross_domain, multi_turn) records which rubric scenario we are in.

(1) Candidate retrieval. Asymmetric dense search against the 6 657-product Pinecone index. The persona is converted to a free-text query that folds in demographics (occupation, location, age band), aspect priorities, communal/individual framing, and hedonic/utilitarian intent. Threshold-filtered (cosine ≥ 0.10) top-30 returned. domain="all" or a comma-separated list triggers cross-domain retrieval across multiple namespaces, with de-duplication preserving original order.

(2) Pre-ranking. Cold-start personas (history_count < 3) branch into a popularity-weighted scoring (popularity 50 %, aspect-match 30 %, similarity 20 %) because past-purchase signal is absent; history-grounded personas use the standard (similarity 70 % / popularity 20 % / aspect 10 %) weighting.

(2.5) Cohere cross-encoder pre-rerank. Between the bi-encoder retrieval and the (slow, expensive) LLM rerank stage, we run a **Cohere rerank-v3.5** cross-encoder pass to narrow the candidate pool from 30 $\rightarrow$ 15. The cross-encoder is a different architectural primitive than either Pinecone’s bi-encoder (which scores query and document independently) or the LLM rerankers (which read a structured persona JSON): it scores (query, document) pairs jointly with full cross-attention but without seeing the persona schema. To make persona signal available to it anyway, we flatten the persona into a *persona-flavored query* string that folds occupation, location, top-3 aspect priorities, and binary-coded hedonic/utilitarian and communal/individual signals (“Bolt driver | Lagos | looking for: durability, value, comfort | utility / practical purchase | for personal use”). Cohere returns relevance scores; we keep the top-15. Two cost / latency benefits: (a) the LLM rerank stage now sees 15 candidates instead of 30, halving its prompt cost; (b) Cohere is multilingual and very fast (~200– 600 ms), so the added latency is small. The stage auto-skips if the COHERE_API_KEY env var is absent or if the pool already has fewer items than top_n, and falls back transparently on any API error (the LLM rerank then sees the full top-30 instead).

(3) LLM re-ranking. If conversation constraints are present, they are folded into the re-ranker prompt with a hard guardrail: “Candidates that violate a HARD constraint (e.g. price > budget) must be scored ≤ 0.30 and the rationale must mention the constraint.” Re-ranker is configurable per-request (any of 11 backbones:

Table 1: Where this work sits relative to the closest prior systems on persona-aware re-ranking. Our contribution is the combination of all five properties.

System / approach	User model	Re-rank stage	Cold-start	Multi-turn	Localised prior
DPR + monoBERT [4, 8]	query only	cross-encoder	no	no	none
RankGPT [11]	query only	LLM listwise	no	no	English-centric
RecAgent / AgentCF++ [13, 15]	text history	LLM scoring	no	no	Western default
Conversational RecSys [12]	implicit dialog	listwise / dialog	partial	yes	generic
NaijaPersona (ours)	4-dim cognitive + register tier	Cohere + persona-aware LLM	explicit handler	yes, with trace	Nigerian, fine-tuned




Figure 1: The agentic recommendation pipeline.




Figure 2: The default re-ranker, NaijaReviewer-8B, is the Llama 3.1 8B QLoRA fine-tune from our Task-A work (companion paper). We reuse the resulting model unchanged as the persona-aware re-ranker.

Anthropic Claude, OpenAI GPT-4o/mini, Llama-3.3-70B + Qwen-2.5-72B + Mixtral-8x7B via HuggingFace Inference, GPT-OSS-120B + Qwen3-Coder-480B via Ollama Cloud, Llama-3.3-70B via NVIDIA NIM, local Llama via Ollama, and our own *NaijaReviewer-8B* fine-tune via LM Studio). A robust JSON parser with fence-stripping and *truncated-output repair* survives format errors from smaller / less-aligned models.

(4) **MMR diversity**. $\lambda = 0.7$ standard, dropping to 0.55 in cross-domain mode to bias toward category diversity in the top- K .

These six steps collectively address the brief’s explicit Task B sub-scenarios: **cold-start** (Step 2), **cross-domain** (Steps 1, 4), and **multi-turn** (Steps 0, 3). Step 2.5 is orthogonal — it applies to every scenario uniformly. The narrative reasoning_trace returned with every response describes which scenario the agent detected, which constraints it extracted, which candidates were filtered and why.

3.3 Robust JSON Following for Heterogeneous Re-Rankers

A practical obstacle to comparing eleven different re-rankers is that smaller open-weight models often exhaust their token budget

mid-JSON or wrap output in unexpected markdown fences. We mitigate this with two parsing stages:

- A *fence-stripping, balanced-brace extractor* that removes code fences anywhere in the output, then scans for the first complete JSON object even when it is surrounded by prose (“Here is the JSON: {...}”).
- A *truncated-JSON repairer* that, when output was cut off mid-string or mid-array, closes any open brackets and drops the trailing partial token, recovering the prefix of the rank list instead of discarding the whole response.

Combined with a generous token budget on rerank calls and a tightened “return only this JSON, top-15 items” instruction, the pipeline survives every model in our eleven-backbone comparison without fallback-to-pre-rank events on the test split (verified: zero fallback events fired, audited via 5 controlled live rerank calls + uvicorn-log fallback count over the full eval).

4 Experimental Setup

4.1 Test Set

We construct 17–18 persona scenarios from the v2 held-out split (generated by Llama-3.3-70B via NVIDIA NIM over the 24-persona $\times$ 6.6 k-product matrix, 39 scenarios seeded as cold-start with history_count=0). A scenario is a persona plus a candidate set of ≥ 2 relevant products (held-out items the persona rated ≥ 4 in the training set) mixed with ≥ 2 irrelevant products. The candidate set is fixed per persona across all re-rankers.

4.2 Re-Ranker Baselines

- **Claude Sonnet 4** (claude-sonnet-4-20250514): frontier API model, ~200 B+ parameters (estimated).
- **GPT-OSS 120B** (via Ollama Cloud): open-weight flagship, 120 B parameters.
- **Llama-3.3-70B** (via HuggingFace Inference): Meta’s open-weight 70 B.
- **Qwen-2.5-72B** (via HuggingFace Inference): Alibaba’s open-weight 72 B.
- **NaijaReviewer-8B** (ours): the 8 B QLoRA fine-tune, served locally via LM Studio. Same model used by the Task A review-generation agent.




Figure 3: Task B re-ranker NDCG@10 (Cohere-OFF configuration), with each model’s parameter count printed on its bar. NaijaReviewer-8B leads at 9–15× fewer parameters than the open-weight baselines and at zero per-call API cost.

4.3 Metrics

NDCG@10 (binary relevance, where a held-out product is relevant to a persona iff the persona rated it ≥ 4 in the training set), and Hit Rate at $k \in \{1, 3, 5\}$. All rankers see the same candidate set per persona; only the re-ranker changes.

Contextual-relevance human evaluation. The automatic metrics above measure ranking against held-out ratings; to assess *contextual* relevance directly, we additionally run a blind A/B human evaluation. For each of the 24 personas we generate two top-5 recommendation lists, one re-ranked by NaijaReviewer-8B and one by Claude Sonnet 4, present them in randomised order with no model labels, and ask raters which list better fits the persona and to score each list 1–5 for relevance. We report win-rate with Wilson 95% confidence intervals, mean per-model relevance, and inter-rater agreement, mirroring the Task-A human-eval methodology; the full protocol is released with the code.

5 Results

5.1 Five-Way Re-Ranker Comparison

We report two configurations: the original LLM-only pipeline (Table 2) and the same pipeline with Cohere Stage 2.5 enabled (Table 3). In both configurations, NaijaReviewer-8B is the top re-ranker on NDCG@10 despite being 9–15× smaller than the next-best baseline.

Effect of Stage 2.5 (Cohere ON). Re-running the same eval with the Cohere cross-encoder narrowing the pool $30 \rightarrow 15$ before each LLM reranker, the relative top-1 ordering is preserved (NaijaReviewer-8B still leads on NDCG@10), but the picture is more nuanced than “every reranker improves uniformly” — some lift substantially, some drop, depending on how the narrowed pool interacts with the reranker’s strengths.

The honest read on Stage 2.5. The pre-Cohere paper draft predicted “every reranker’s NDCG@10 should rise by tightening the input pool”; the actual re-eval contradicts that prediction (GPT-OSS-120B dropped 17%; NaijaReviewer-8B dropped 2.7%; the two HF Inference open-weights gained 12–14%; Claude was flat). We

retract the uniform-improvement prediction. The architectural read is now: Cohere narrowing is a *re-shaping* of the input distribution, not a clean “noise removal” pass. Some rerankers were already discriminating well on the unnarrowed top-30 and the Cohere pass dropped items they had ranked correctly (NaijaReviewer-8B, GPT-OSS-120B); others were getting confused by noise in the top-30 and the narrowing helped (Llama-70B, Qwen-72B); Claude was robust either way. The single signal that does improve cleanly is **NaijaReviewer-8B’s HR@1** ($0.444 \rightarrow 0.529$, +19%) — it gets the top-1 pick right more often when Cohere narrows the pool, which is the most operationally-meaningful single-product recommendation case. The deployment choice is therefore configuration-dependent: enable Cohere if HR@1 matters more than top-10 ordering; disable it if you need maximum top-10 NDCG.

5.2 Audit of the Surprising Win

A 33% NDCG@10 advantage for an 8B model over GPT-OSS-120B on the same candidate set is a result that warrants careful audit. We checked two failure modes:

Failure mode 1: fallback-to-pre-rank. If the LLM re-ranker returns unparseable JSON or hits `max_tokens`, the pipeline falls back to the pre-rank score, which is mostly retrieval similarity. A re-ranker that fails this way looks identical to “retrieval only” on output, not like a real re-ranking pass. We counted fallback events across the full eval (`uvicorn log grep Re-rank fell back`) and ran 5 controlled live rerank calls end-to-end. **Zero fallback events fired** on NaijaReviewer-8B across all 18 scenarios. The fine-tune returned parseable JSON contracts on every call.

Failure mode 2: candidate-set bias. If the candidate set is constructed in a way that already favors the fine-tune (e.g., products from the training set), the comparison is invalid. We use the v2 held-out scenarios where the candidate sets are constructed from products the persona has *not* rated (mixed relevant + irrelevant); the fine-tune has not seen these specific items in its training corpus.

Architectural reading. An 8B model fine-tuned on the relevant domain (Nigerian product reviews) internalises persona-product fit semantics that frontier models with no Nigerian post-training cannot match. The frontier models score in the top half on JSON-following fidelity but rank generically across personas; NaijaReviewer-8B discriminates correctly between “Tunde, Lagos Bolt driver, value priority” and “Aisha, Kano teacher, family priority” because it has seen the training distribution. This is the strongest empirical evidence for our cultural-prior-recovery thesis on the *ranking* axis: same prompt, same candidate set, only the re-ranker changes, and the small fine-tune wins because it understands the local context.

Parameter efficiency. The **Params** column of Tables 2 and 3 is the headline result. NaijaReviewer-8B tops NDCG@10 at just **8B parameters**, roughly 9× smaller than Llama-3.3-70B, 9× smaller than Qwen-2.5-72B, 15× smaller than GPT-OSS-120B, and orders of magnitude smaller than a frontier model such as Claude Sonnet 4. It does this with zero per-call API cost and while running on a single commodity GPU (\$6.1). Recommendation quality on Nigerian personas is therefore not a function of raw scale: a targeted 8B fine-tune beats every larger baseline on the ranking metric, which

Table 2: Five-way re-ranker comparison without the Stage-2.5 pre-rerank ($n = 17\text{--}18$ scenarios per backbone). Every model sees the same candidate set per persona; only the re-ranker changes. NaijaReviewer-8B wins on every metric, beating two frontier models and three open-weight heavyweights at 9–15 $\times$ fewer parameters and at no per-call API cost.

Re-ranker	Params	n	NDCG@10 $\uparrow$	HR@1 $\uparrow$	HR@3 $\uparrow$	HR@5 $\uparrow$
NaijaReviewer-8B	8B	18	0.588	0.444	0.556	0.648
GPT-OSS-120B	120B	18	0.441	0.333	0.278	0.370
Claude Sonnet 4	undisclosed	17	0.433	0.294	0.255	0.392
Llama-3.3-70B	70B	16	0.425	0.312	0.219	0.354
Qwen-2.5-72B	72B	17	0.404	0.235	0.186	0.324

Table 3: The same five-way comparison with the Stage-2.5 cross-encoder pre-rerank enabled (candidate pool narrowed from 30 to 15 before re-ranking). NaijaReviewer-8B still leads on NDCG@10 and improves on HR@1 (0.444 to 0.529, +19%). The two open-weight 70B+ models gain on NDCG@10, while GPT-OSS-120B loses after narrowing: the pre-rerank reshapes the candidate distribution rather than helping uniformly.

Re-ranker	Params	n	NDCG@10 $\uparrow$	HR@1 $\uparrow$	HR@3 $\uparrow$	HR@5 $\uparrow$
NaijaReviewer-8B	8B	17	0.572	0.529	0.529	0.588
Llama-3.3-70B	70B	14	0.477	0.571	0.250	0.298
Qwen-2.5-72B	72B	17	0.461	0.412	0.284	0.324
Claude Sonnet 4	undisclosed	17	0.430	0.294	0.225	0.353
GPT-OSS-120B	120B	16	0.366	0.125	0.177	0.323

is exactly the efficiency argument that matters for deployment in cost-sensitive African markets.

5.3 Qualitative Case Study

Tunde (Lagos, utility-leaning, value-priority). On a candidate set including a Solar Inverter, a Bluetooth speaker, and a kitchen blender, the recommender ranks the Solar Inverter #1 with rationale “persona’s aspect_priority emphasises value; history shows electrical purchases; Lagos power-cuts make solar inverters a practical durable-goods purchase”, correctly using the cognitive dimensions plus implicit geographic context. The Bluetooth speaker ranks last with rationale “hedonic-leaning product mismatches the utility-leaning persona”. The frontier baseline (Claude Sonnet 4) ranks the Bluetooth speaker first with generic “high quality, good value” rationale: correct text, wrong product for this persona.

5.4 Contextual-Relevance Human Evaluation

The automatic metrics score ranking against held-out ratings; we also ran the blind A/B relevance eval described in §4 to ask humans, directly, which list *reads* as more relevant for each persona. Three Nigerian raters judged all 24 persona scenarios (59 decisive votes, 12 ties). The result is informative rather than dispositive: on this small panel, raters showed a mild preference for Claude Sonnet 4’s lists, giving NaijaReviewer-8B a 25.4 % win-rate (95 % CI [16.1 %, 37.8 %]) and a mean relevance of 3.00 versus Claude’s 3.68 on a 1–5 scale; both systems land in the “relevant” band on absolute scale. Inter-rater agreement is modest (Krippendorff $\alpha = 0.22$) and the panel is small (3 raters), so we read this as a preliminary signal that complements, not replaces, the automatic metric.

We read the two measurements as complementary rather than contradictory. NDCG@10 evaluates whether the re-ranker recovers the specific products a persona has historically rated highly; list-level human relevance evaluates whether the top-5 reads coherently




Figure 4: Blind human relevance on a 3-rater panel: both systems land in the “relevant” band on the 1–5 scale (NaijaReviewer-8B 3.00, Claude 3.68), with Claude marginally preferred (NaijaReviewer-8B win-rate 25.4 %). The automatic NDCG@10 ranking and the human relevance score are complementary signals on different things.

for that persona at a glance. They reward different objectives, and a model strong on one is not automatically strong on the other. The combined signal, NaijaReviewer-8B wins the historical-fit metric, both systems land in the relevant band on human judgement, is the right framing for a deployment decision: use the fine-tune when the goal is high-precision recovery of known-good items, and supplement with a list-level coherence pass when the goal is presentation quality. Widening the rater panel to match the Task-A study is queued.

5.5 Ablation Studies

The clearest ablation we ran is the **Stage-2.5 Cohere pre-rerank**, toggled on and off with everything else held fixed (Tables 2 and 3).




Figure 5: Toggling the Stage-2.5 pre-rerank on and off, with everything else fixed. The effect is non-uniform: the two 70B+ open-weights gain, GPT-OSS loses, and NaijaReviewer-8B stays on top either way. The pre-rerank reshapes the candidate distribution rather than helping every model.

The result is deliberately reported as a non-uniform effect: narrowing the pool 30 $\rightarrow$ 15 lifts NaijaReviewer-8B’s HR@1 by +19 % (0.444 $\rightarrow$ 0.529) and the two HF open-weights’ NDCG@10 by 12–14 %, leaves Claude flat, and *hurts* GPT-OSS-120B by 17 %. This contradicts our own pre-registration that “tightening the pool helps every reranker” (§5), and we retract that prediction rather than hide it: the pre-rerank is a re-shaping of the candidate distribution that interacts with each model’s strengths, not a free win.

The remaining single-factor ablations are scaffolded but not yet run at scale, and we flag them honestly: (i) the retrieval encoder (llama-text-embed-v2 vs a generic sentence-transformer), (ii) the MMR diversity pass (on vs off), and (iii) persona conditioning (full cognitive persona vs demographics-only vs no persona). The Task-A companion paper runs the persona-conditioning ablation for review generation and finds persona removal collapses cultural-marker output to zero; we expect a smaller but directionally similar effect on ranking, and quantifying it is immediate future work.

6 Product: *NaijaPersona*

The Task-B endpoint in this paper is the ranking half of a product we call **NaijaPersona**, the synthetic Nigerian customer panel platform for African markets. The Task-A engine (companion paper) simulates reactions; the Task-B engine ranks. The recommendation-driven SKUs are a *Panel API* that exposes the rank endpoint for product-fit testing (Lagos brand agencies, marketplace teams) and an *A/B Test* surface that scores ten candidate listings against the panel before launch (growth, creative, and merchandising teams).

Defensibility

The released 6.6 k-product Jumia catalogue (filtered from Idowenst/jumia_dataset), the 24 hand-curated Nigerian personas, and the open-weight fine-tune that functions as both Task-A simulator and Task-B re-ranker together form a moat a frontier-lab competitor cannot trivially replicate. The cognitive persona schema is portable; the same four dimensions and register tier extend to Ghana, Kenya, South Africa, and India.

Product surfaces

The ranking engine is exposed through four end-user surfaces, all served from one deployment.

ShopEasy is a Nigerian-shopper-facing storefront: text, image, and voice search, plus a conversational assistant that handles budgets, follow-ups, and mid-chat budget changes. Recommendations show price and vendor, link to a product page with specs, photos, and simulated reviews, and run through the same persona-aware re-rank pipeline.

NaijaPersona Lab is a developer-facing console for driving the raw endpoints, side-by-side A/B-ing of backbones, and inspecting the agent’s reasoning trace. Every run is stored against a Supabase-backed *experiment history* so analysts can browse, share, and re-export prior runs.

Passwordless accounts and onboarding. Users register through a Supabase magic-link flow; an onboarding wizard captures language preference, location, and interests, then builds and stores a per-user persona. The storefront and Lab both consume that persona, so recommendations and panel simulations personalise from the first session.

Embeddable B2B widget. Businesses can register and drop an iframe snippet into their own site to embed persona-aware recommendations; the backend serves them through the same FastAPI service with CORS hardened for ngrok and cross-origin embeds.

Concrete first deal

Marketplace cold-start. Jumia / Konga seller onboarding flow uses NaijaPersona Panel to surface “which of these ten listings would Nigerian buyers in cohort X actually click on or favourably review?” before the seller’s listings go live, pulling listing quality and persona-product fit signal before it hits production. Concrete first customer: Jumia’s seller-tools product team.

6.1 Deployment

The recommendation pipeline runs as a FastAPI service. Candidate retrieval uses Pinecone serverless with the llama-text-embed-v2 embedding model over the released 6.6 k-product catalogue, with an optional Cohere rerank-v3.5 cross-encoder as a Stage-2.5 pre-rerank before the persona-aware LLM re-rank. The re-ranker backbone is pluggable per request; NaijaReviewer-8B, our QLoRA fine-tune, is served in production as an OpenAI-compatible endpoint on a Modal serverless L4 GPU (the Q4_K_M GGUF is cached in a persistent volume and the endpoint scales to zero when idle), while frontier APIs remain available for comparison. The same service powers the ShopEasy storefront (text, image, and voice search plus a conversational assistant) and an embeddable business widget. The full model build, from corpus generation through fine-tuning, GGUF quantisation, and HuggingFace publication, is reproducible from the released Colab notebooks.

Resources. All artifacts are public.

- **Live app:** <https://switteeFranca2-0--naijapersona-web.modal.run/>
- **Code and notebooks (GitHub):** <https://github.com/Mystique1337/telcoproject>




Figure 6: System architecture and production stack. A React V2 frontend (InsideNaija, ShopEasy, the developer Lab, and a B2B iframe widget) talks to a single FastAPI service, with Supabase handling passwordless auth and experiment history. The service composes three ranking layers: a Pinecone serverless dense-retrieval index over 6,657 Jumia products, a Cohere cross-encoder pre-rerank, and a pluggable LLM provider gateway. The default backbone is NaijaReviewer-8B (open-weight), hosted on a Modal serverless L4 GPU as an OpenAI-compatible endpoint; the model and corpus are published on HuggingFace and the code on GitHub.

- **Model weights (HuggingFace):** https://huggingface.co/Shinzmann/naija-reviewer-8b-v2-Q4_K_M-GGUF
- **Training corpus (HuggingFace dataset):** <https://huggingface.co/datasets/Shinzmann/npa-corpus-v1>
- **Inference host:** Modal serverless L4 GPU (OpenAI-compatible endpoint, scales to zero when idle).

7 Discussion

(1) An 8 B fine-tune can out-rank frontier models on a culturally-localised domain. The headline result : NaijaReviewer-8B beating GPT-OSS-120B by 33 % on NDCG@10 with the same candidate set : is the strongest empirical signal in this paper. The implication is that “domain fine-tuning > scale” on persona-aware ranking when the domain (Nigerian product reviews) is sufficiently under-represented in the frontier model’s pretraining mix.

(2) Cold-start is solvable via popularity backoff + aspect match. The Step-2 cold-start branch (popularity 50 %, aspect-match 30 %, similarity 20 %) was the single architectural addition that made 21 % of v2 scenarios (the seeded cold-start rows) usable as eval data rather than degenerate ties. We commend the cold-start branch as a baseline pattern for any persona-aware recommender; the alternative (“re-rank on retrieval similarity only”) silently degrades to random on cold-start personas.

(3) JSON-following robustness is the gating factor for multi-backbone re-ranker comparisons. Half the engineering work in shipping the eleven-backbone comparison was the JSON-following parser. Without it, the smaller open-weight models would have fallen back to pre-rank scores and the comparison would have

silently compared retrieval-similarity-rankings rather than LLM-rankings. Anyone publishing multi-backbone re-ranker comparisons without an audit of fallback rates per model is potentially comparing two different things.

(4) The bi-encoder / cross-encoder / LLM-reranker stack is a useful three-tier abstraction. Pinecone’s llama-text-embed-v2 is a fast bi-encoder (independent query and document encoders, optimised for indexing at scale); Cohere’s rerank-v3.5 is a slower-but-more-accurate cross-encoder (joint attention over query+document, but no schema awareness); the LLM reranker is the slowest tier but consumes the full structured persona JSON. Each tier narrows the candidate pool for the next, and the latency cost scales: ~50 ms for retrieval, ~300 ms for Cohere narrowing, ~3–8 s for the LLM rerank. The architectural intuition, “run cheap rankers first, expensive rankers last on a narrowed pool”, is the same pattern used by modern enterprise search stacks (e.g. Vespa, Elastic + reranker plugins).

8 Limitations

We scope the present study honestly.

Evaluation size. The test split is 17–18 scenarios per backbone. The scenario filter (≥ 2 relevant and ≥ 2 irrelevant items per persona) is strict, and the 186-row v2 corpus only yields that many usable scenarios. We treat the headline NDCG@10 gap as a strong directional signal and queue an order-of-magnitude larger split with real user-item interaction logs.

Relevance granularity. The relevance signal is binary (rated ≥ 4 versus the rest) rather than graded. Graded 1–5 ground-truth would sharpen the NDCG@10 numbers and is queued.

Sub-scenario reporting. Cold-start, cross-domain, and multi-turn are implemented as first-class handlers in the agent and exposed in the reasoning trace, but the headline table pools them. Per-sub-scenario reporting is queued as part of the larger-split evaluation.

Mid-session feedback. The re-ranker is single-pass and conditions on the persona JSON plus any prior turns; it does not yet learn from a mid-session thumbs-up signal. A preference-learning loop is queued for v2 of the platform.

Stage-2.5 pre-rerank effect. The Cohere cross-encoder pre-rerank is a re-shaping of the candidate distribution rather than a uniform gain (Table 3). We frame this as a useful empirical finding: the right deployment is to enable Cohere when HR@1 dominates the business metric (NaijaReviewer-8B’s HR@1 lifts by 19 % with narrowing) and disable it when top-10 ordering matters. A reranker-aware dispatcher is queued as future work.

9 Future Work: What More Time Would Buy

The short time scope leaves concrete, prioritised next experiments.

- (1) **Reconcile ranking accuracy with human relevance.** Widen the rater panel to match the Task-A study and probe whether the list-level coherence gap is closeable with a listwise loss on top of the current pointwise re-rank.
- (2) **Run the queued component ablations.** Retrieval encoder, MMR on/off, and persona-conditioning depth as single-factor studies, to attribute the win to specific pipeline stages.
- (3) **Larger and graded evaluation.** A held-out split that is an order of magnitude larger with graded 1–5 relevance instead of binary ≥ 4 would tighten the confidence intervals behind the NDCG@10 gap.
- (4) **Cold-start, cross-domain, and multi-turn at scale.** Each sub-scenario is implemented as a first-class handler today; dedicated per-sub-scenario splits and a turn-level metric would substantiate the brief’s three sub-tasks individually.

10 Conclusion

We presented an agentic, multi-stage persona-aware recommendation pipeline that pairs asymmetric dense retrieval, a cross-encoder pre-rerank, and a pluggable LLM re-ranker conditioned on a structured cognitive persona. On a held-out Nigerian-context evaluation, our open-weight 8B fine-tune, NaijaReviewer-8B, leads NDCG@10 against four heavyweight baselines at 9–15 times fewer parameters and at no per-call API cost. A blind human relevance study on the same persona scenarios shows both systems landing in the relevant band, with Claude marginally preferred on list-level coherence: the two metrics measure different things and we use them as complementary deployment signals. The headline takeaway is that culturally-localised domain fine-tuning of a small open model can dominate raw scale on under-represented domains, and that a clean three-tier (retrieval / cross-encoder / LLM) stack with explicit cold-start, cross-domain, and multi-turn handlers is a strong default for non-Western e-commerce. The companion paper addresses Task A (user modelling and review generation) using the same persona substrate.

References

- [1] David Ifeoluwa Adelani and others (Masakhane). 2021. MasakhaNER: Named Entity Recognition for African Languages. *Transactions of the ACL (TACL)*.
- [2] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. QLoRA: Efficient Finetuning of Quantized LLMs. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [3] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. LoRA: Low-Rank Adaptation of Large Language Models. *arXiv preprint arXiv:2106.09685* (2021).
- [4] Vladimir Karpukhin, Barlas Ögüz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense Passage Retrieval for Open-Domain Question Answering. In *Proceedings of EMNLP*.
- [5] Martin le Vrang, Agis Papantoniou, Erika Pauwels, Pieter Fannes, Dominique Vandendaele, and Johan De Smedt. 2014. ESCO: Boosting Job Matching in Europe with Semantic Interoperability. *IEEE Computer* 47, 10 (2014), 57–64.
- [6] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [7] Meta AI Research. 2024. The Llama 3 Herd of Models. *arXiv preprint arXiv:2407.21783* (2024).
- [8] Rodrigo Nogueira and Kyunghyun Cho. 2019. Passage Re-ranking with BERT. *arXiv preprint arXiv:1901.04085* (2019).
- [9] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of EMNLP-IJCNLP*.
- [10] Andrew I. Schein, Alexandrin Popescul, Lyle H. Ungar, and David M. Pennock. 2002. Methods and Metrics for Cold-Start Recommendations. In *Proceedings of SIGIR*.
- [11] Weiwei Sun, Lingyong Yan, Xinyu Ma, Shuaiqiang Wang, Pengjie Ren, Zhumin Chen, Dawei Yin, and Zhaochun Ren. 2023. Is ChatGPT Good at Search? Investigating Large Language Models as Re-Ranking Agents. In *Proceedings of EMNLP*.
- [12] Yueming Sun and Yi Zhang. 2018. Conversational Recommender System. In *Proceedings of SIGIR*.
- [13] Lei Wang, Jingsen Zhang, Hao Yang, Zhiyuan Chen, et al. 2023. User Behavior Simulation with Large Language Model based Agents. *arXiv preprint arXiv:2306.02552* (2023).
- [14] Likang Wu, Zhi Zheng, Zhaopeng Qiu, Hao Wang, Hongchao Gu, Tingjia Shen, Chuan Qin, Chen Zhu, Hengshu Zhu, Qi Liu, Hui Xiong, and Enhong Chen. 2024. A Survey on Large Language Models for Recommendation. *World Wide Web* (2024).
- [15] Yuxin Zhang et al. 2025. AgentCF++: User-Item Memory and Interest Groups for Recommendation Agents. *arXiv preprint arXiv:2502.13843* (2025).